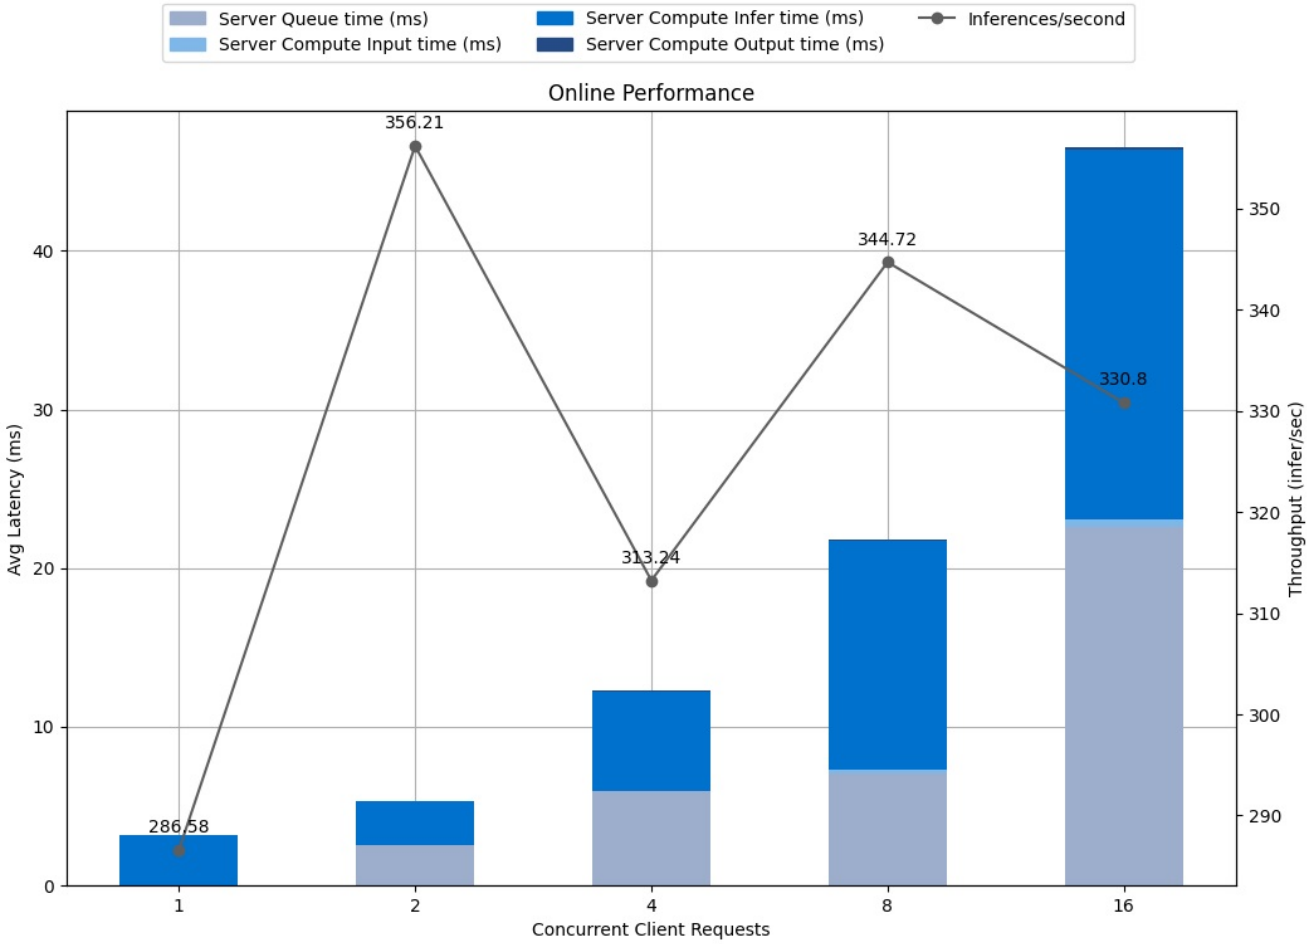
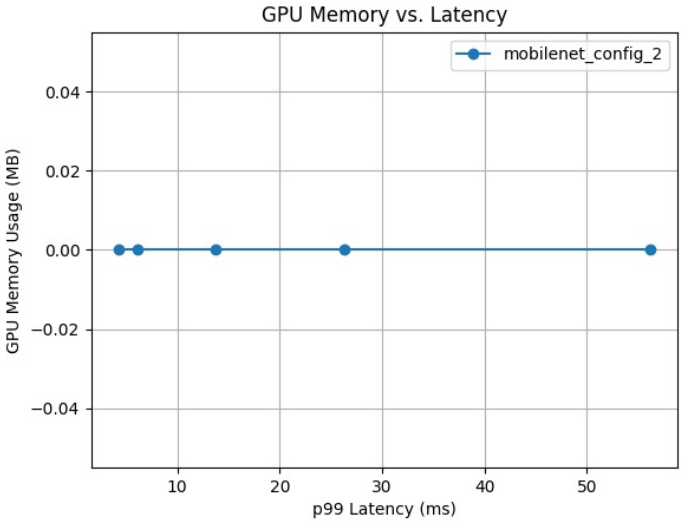
Detailed Report

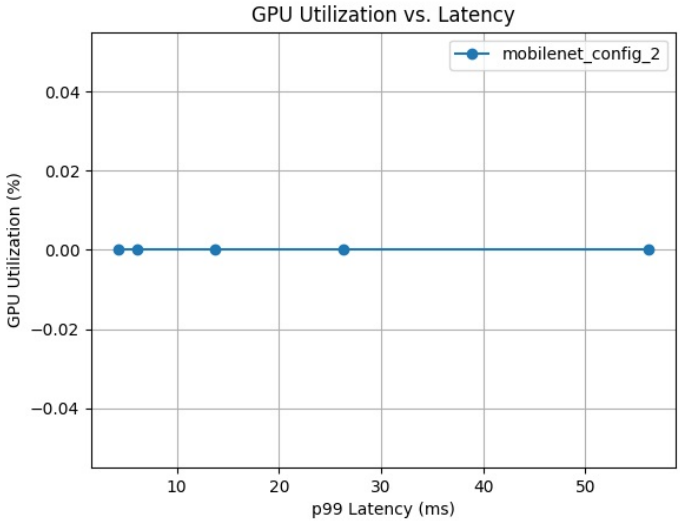
Model Config: mobilenet_config_2



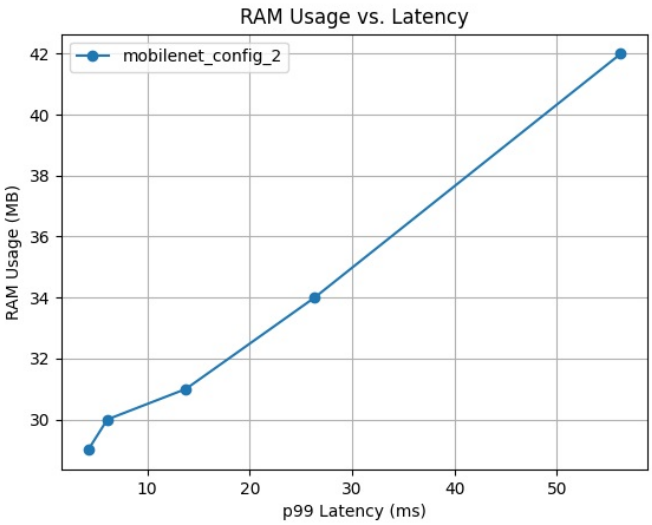
Latency Breakdown for Online Performance of mobilenet_config_2



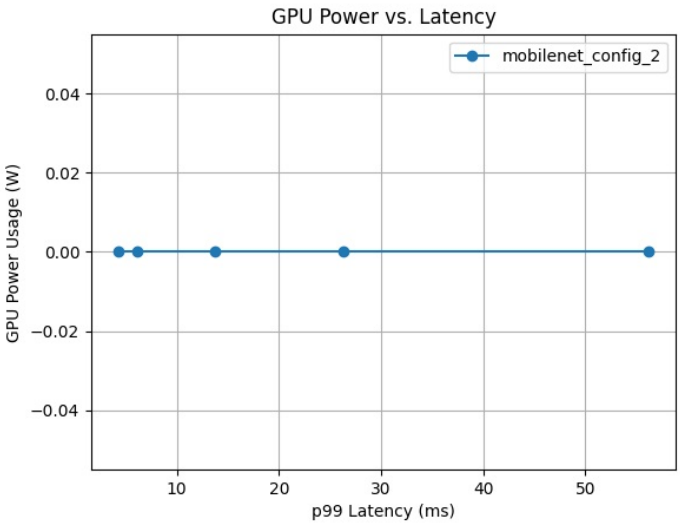
GPU Memory vs. Latency curves for config mobilenet_config_2



GPU Utilization vs. Latency curves for config mobilenet_config_2



RAM Usage vs. Latency curves for config mobilenet_config_2



GPU Power vs. Latency curves for config mobilenet_config_2

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max CPU Memory Usage (MB)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
16	56.259	47.864	22.602	0.435	23.326	330.796	42.0	0	0
8	26.348	23.032	7.062	0.244	14.379	344.719	34.0	0	0
4	13.761	12.661	5.92	0.055	6.26	313.241	31.0	0	0
2	6.094	5.567	2.537	0.024	2.754	356.21	30.0	0	0
1	4.213	3.421	0.02	0.017	3.137	286.577	29.0	0	0

The model config **mobilenet_config_2** uses 1 CPU instance with a max batch size of 8 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.