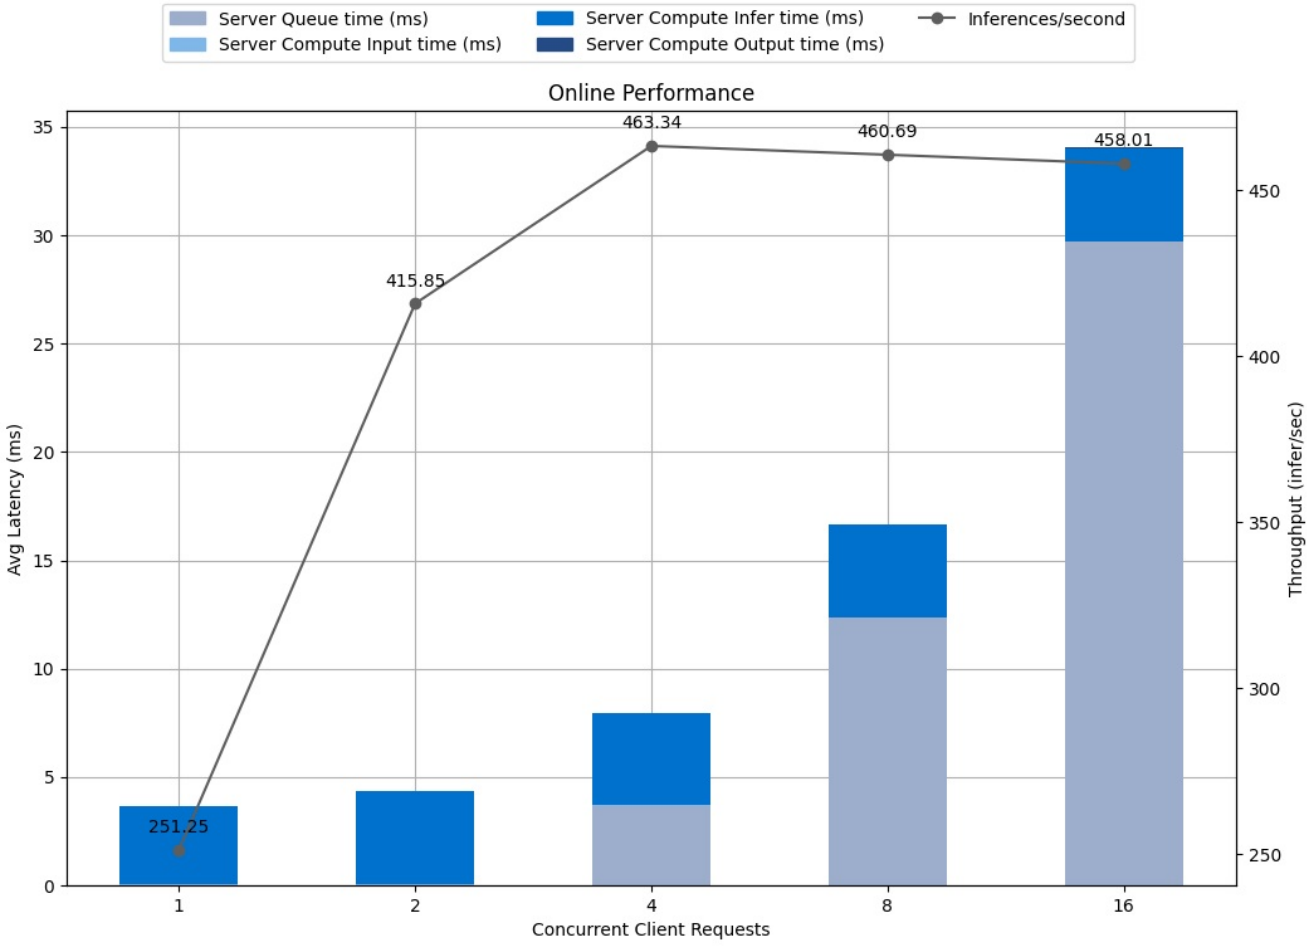
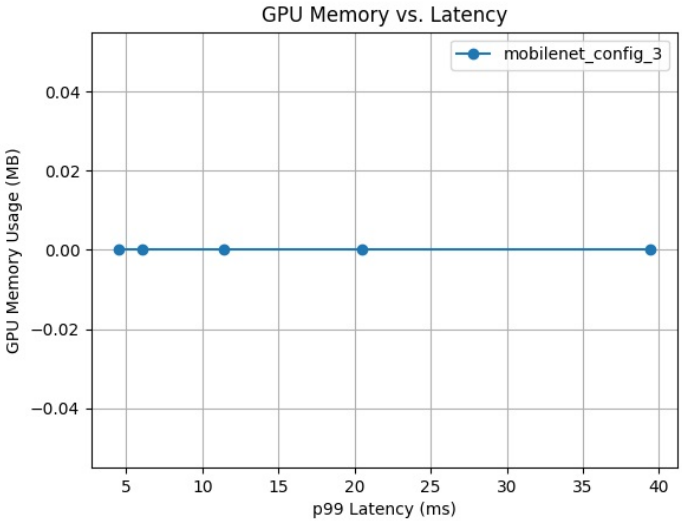
Detailed Report

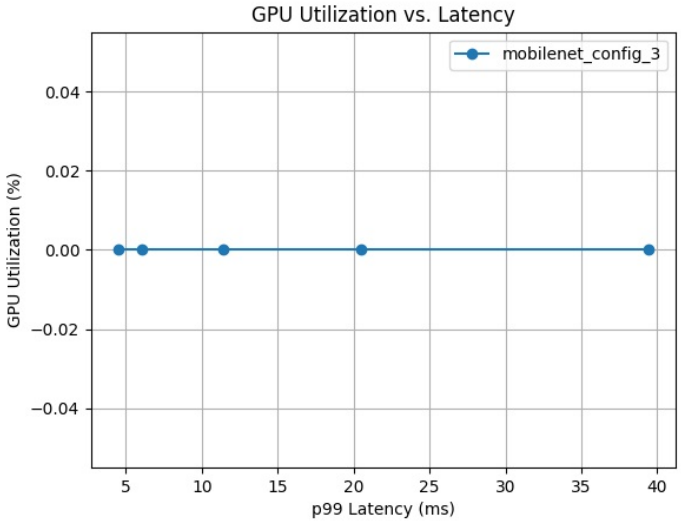
Model Config: mobilenet_config_3



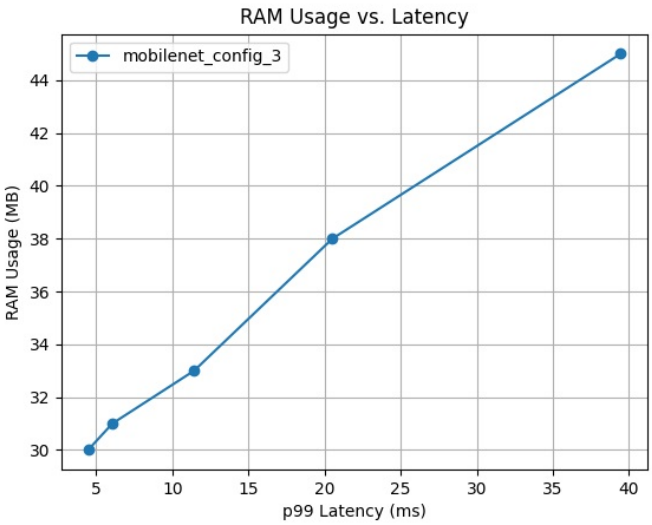
Latency Breakdown for Online Performance of mobilenet_config_3



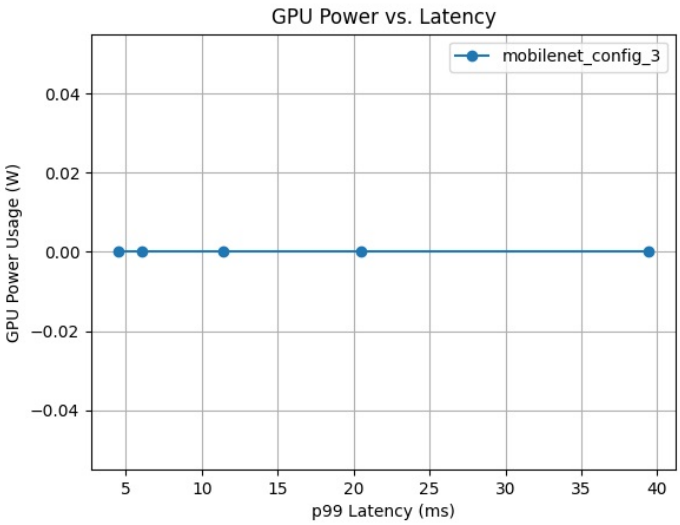
GPU Memory vs. Latency curves for config mobilenet_config_3



GPU Utilization vs. Latency curves for config mobilenet_config_3



RAM Usage vs. Latency curves for config mobilenet_config_3



GPU Power vs. Latency curves for config mobilenet_config_3

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max CPU Memory Usage (MB)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
16	39.457	34.582	29.705	0.039	4.28	458.009	45.0	0	0
8	20.524	17.214	12.345	0.04	4.249	460.692	38.0	0	0
4	11.437	8.534	3.666	0.037	4.232	463.344	33.0	0	0
2	6.068	4.716	0.027	0.029	4.297	415.85	31.0	0	0
1	4.474	3.918	0.021	0.02	3.624	251.254	30.0	0	0

The model config **mobilenet_config_3** uses 2 CPU instances with a max batch size of 1 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.