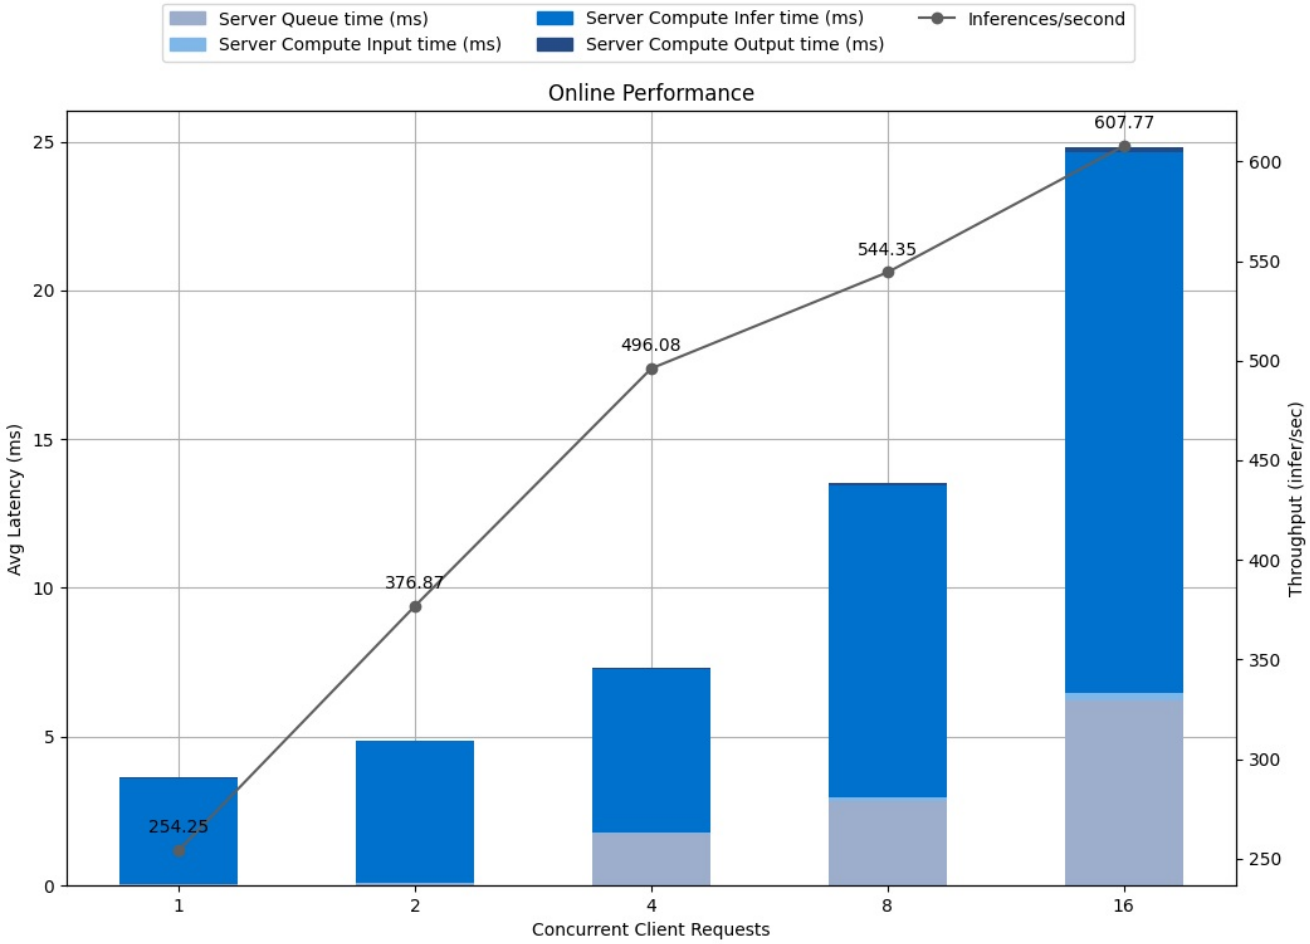
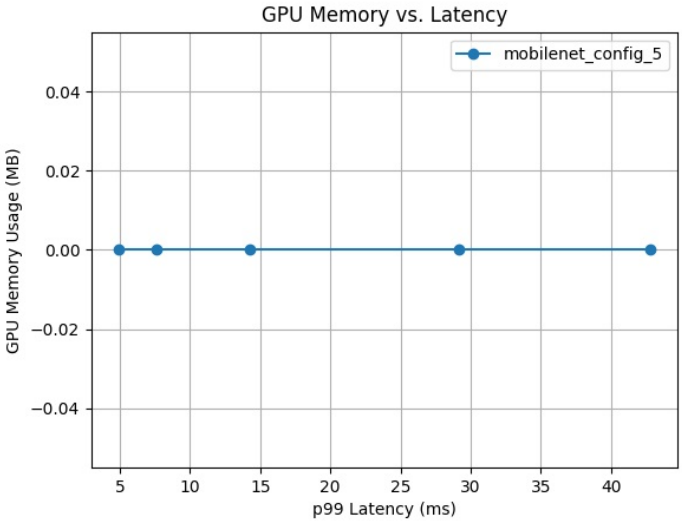
Detailed Report

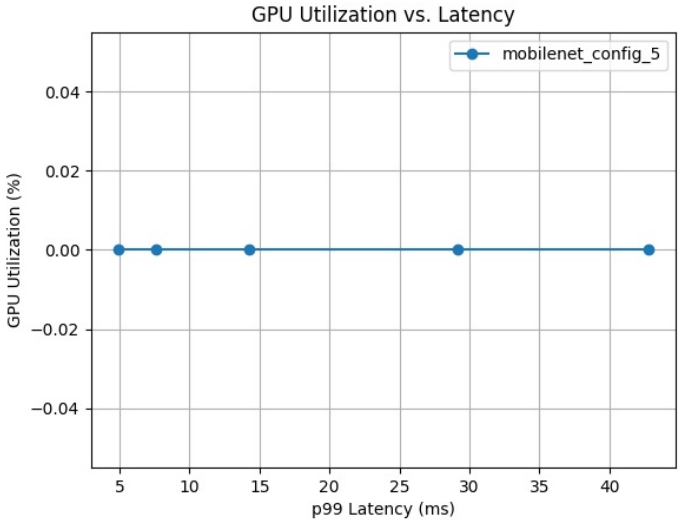
Model Config: mobilenet_config_5



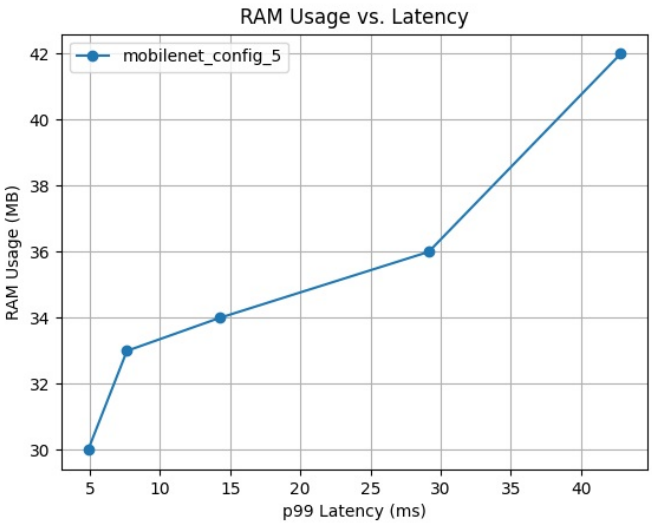
Latency Breakdown for Online Performance of mobilenet_config_5



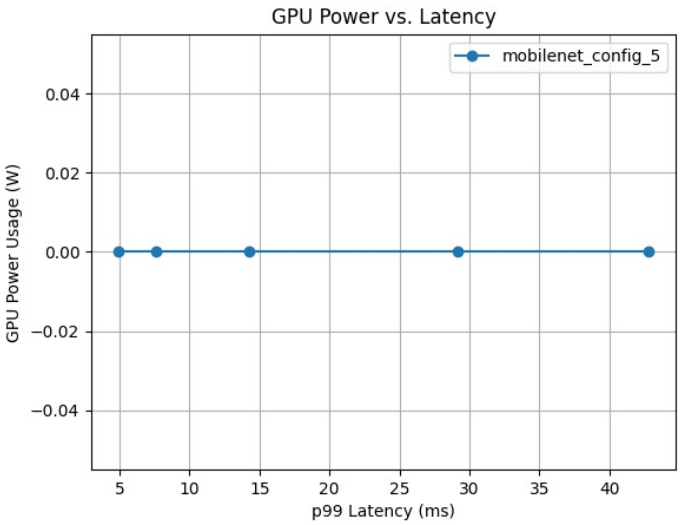
GPU Memory vs. Latency curves for config mobilenet_config_5



GPU Utilization vs. Latency curves for config mobilenet_config_5



RAM Usage vs. Latency curves for config mobilenet_config_5



GPU Power vs. Latency curves for config mobilenet_config_5

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max CPU Memory Usage (MB)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
16	42.825	26.034	6.208	0.266	18.184	607.773	42.0	0	0
8	29.144	14.6	2.839	0.136	10.481	544.349	36.0	0	0
4	14.261	7.952	1.74	0.048	5.502	496.081	34.0	0	0
2	7.647	5.21	0.029	0.036	4.78	376.874	33.0	0	0
1	4.884	3.863	0.022	0.021	3.562	254.251	30.0	0	0

The model config **mobilenet_config_5** uses 2 CPU instances with a max batch size of 8 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.